



arm

Armv8-M Mainline Floating-point Extension

Learning Objectives

At the end of this topic, you will be able to:

- Summarize the Floating-point Extension features in the Armv8-M architecture
- Describe the Floating-point register bank and Floating-point system control registers
- Enable the Floating-point Extension from software
- Describe the mnemonic structure of Floating-point instructions
- Define the types of Floating-point exceptions

Agenda

- **Floating-point Extension overview**
- Registers
- Enabling the FPU
- Floating-point instructions
- Exceptions

Floating-point extension overview

- **The optional Floating-point Extension defines a Floating Point Unit (FPU)**
 - Only available in Mainline
- **Version of Floating-point Extension supported is FPv5**
- **IEEE 754-2008 compliant**
- **Provides:**
 - Single-precision arithmetic
 - Optional double-precision arithmetic (synthesis option)
 - 32 single-precision Floating-point registers S0-S31, also viewable as 16 double precision registers D0-D15
 - Data transfers between core registers and Floating-point registers of single-precision and double-precision values
 - Conversions between integer, half-precision, single-precision and double-precision formats
 - Software selectable Flush-to-zero mode

Floating-point extension overview

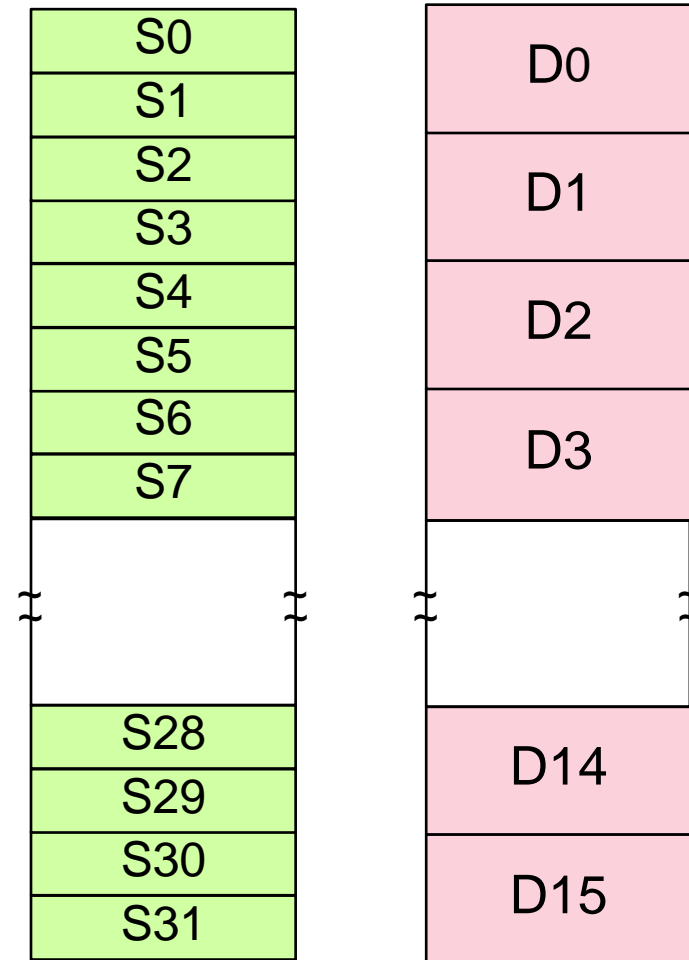
- **Encoding of instructions are for coprocessors 10 and 11**
 - Historical reasons only
 - Fully integrated into processor core
- **Adds system registers in coprocessor system space and System Control Block (SCB)**
 - System register in coprocessor system space – FPSCR
 - System registers in System Control Space – FPCAR, FPCCR, FPDSCR, MVFR0, MVFR1, MVFR2
- **Software can interrogate Media and VFP Registers (MVFR0, MVFR1, MVFR2) to discover the implemented Floating-point features**

Agenda

- Floating-point Extension overview
- **Registers**
- Enabling the FPU
- Floating-point instructions
- Exceptions

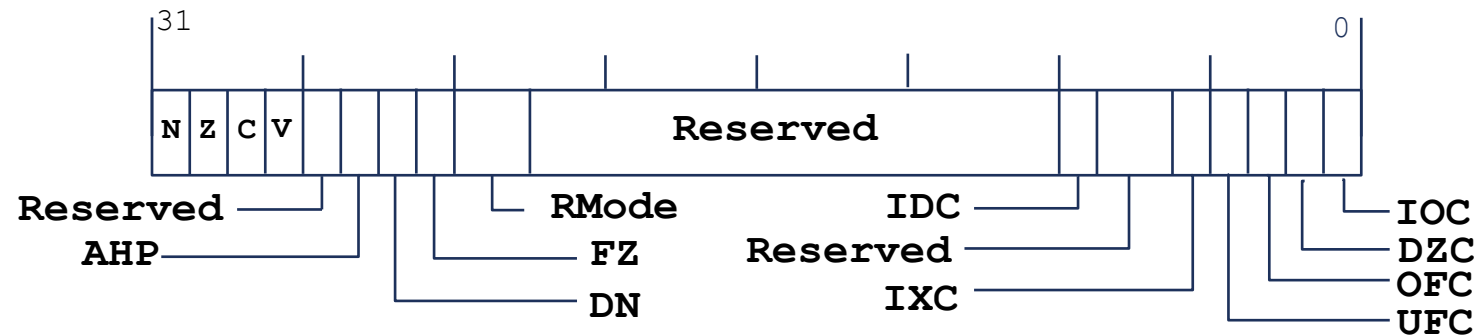
Data-processing Floating-point registers

- FPU provides a 32 single-precision registers
- Can be viewed as either
 - 32 x 32-bit registers
 - 16 x 64-bit doubleword registers
 - Any combination of the above
- Not banked between security states



Floating-point Status and Control Register


- Floating-point Status and Control Register (FPSCR)
- Privileged access only, unprivileged access is RAZ/WI
- Not banked between security states
- Value is UNKNOWN on reset



- **N, Z, C, V** – Floating-point comparison flags
- **AHP** – Alternative (non-IEEE 754-2008) half-precision
- **DN** – default NaN
- **FZ**– Flush-to-zero mode, any denormal inputs or results are flushed to zero (non-IEEE 754-2008)
- **Rounding Mode** Round to Nearest, Round towards Plus Infinity, Round towards Minus Infinity, Round towards Zero
- **IDC, IXC, IOC, DZC, OFC, UFC** – cumulative exception bits

Floating-point Default Status Control Register

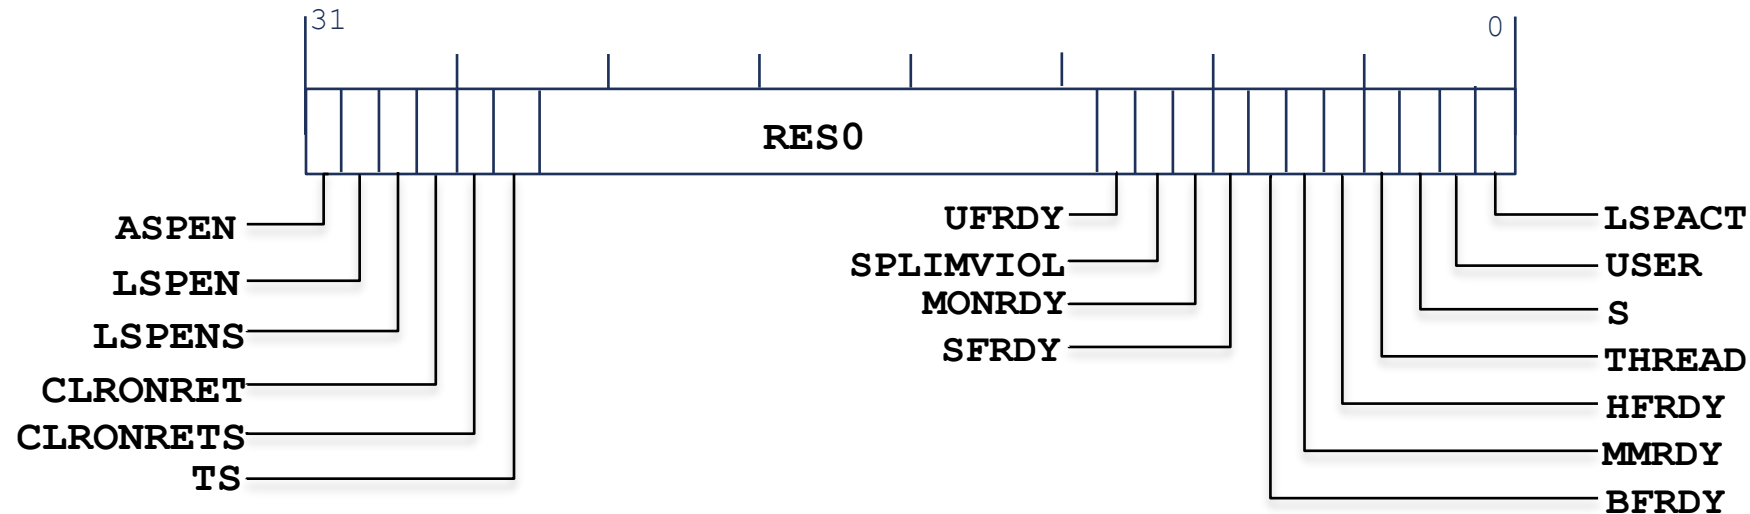
- Floating-point Default Status Control Register (FPDSCR)
- Holds default values for certain fields of the FPSCR when creating a new Floating-point context
- Privileged access only, unprivileged access generates a fault
- Banked between security states (non-secure version is FPDSCR_NS)
- All non-reserved fields set to zero on warm reset



- AHP – Alternative half-precision
- DN – Default NaN
- FZ – Flush-to-zero mode
- Rmode – Rounding mode

Floating-point Context Control Register

- Floating-point Context Control Register (FPCCR)
- Holds control data for the floating-point unit
- Privileged access only, unprivileged access generates a fault
- Banking between security states is on a per-bit basis
- Non-secure versions is FPCCR_NS, not all bits present in non-secure version



Agenda

- Floating-point Extension overview
- Registers
- **Enabling the FPU**
- Floating-point instructions
- Exceptions

Enabling the FPU in CPACR

- Until the FPU is enabled, any Floating-point instruction or access to a register is **UNDEFINED**, causing a UsageFault exception
- To enable access to the FPU, the access bits [23:20] for CP10 and CP11 have to be enabled
- Coprocessor Access Control Register CPACR for secure state
 - 0b00 == no access
 - 0b01 == privileged access only
 - 0b10 == **RESERVED**
 - 0b11 == any access
- **CP10 and CP11 access have to be set the same, otherwise result is UNPREDICTABLE**

```
LDR    R0, =0xE000ED88    ; Load address of CPACR
LDR    R1, [R0]           ; Load current CPACR
ORR    R1, R1, #(0xF << 20) ; Set CP10, CP11 any access
STR    R1, [R0]           ; Write modified CPACR value
ISB                               ; Instruction Synchronization Barrier
```

Enabling the FPU for Non-secure state

- **Non-Secure Access Control Register should be programmed to enable FPU in Non-secure state along with CPACR settings**
- **Non-secure Access Control Register NSACR for non-secure state**
 - 0b0 == non-secure accesses generate NOCP fault
 - 0b1 == any access is allowed
- **CP10 and CP11 access have to be set the same, otherwise result is UNPREDICTABLE**

```
LDR    R0, =0xE000ED8C      ; Load address of NSACR
LDR    R1, [R0]              ; Load current NSACR
ORR    R1, R1, #(0x3 << 10) ; Set CP10, CP11 non-secure access allowed
STR    R1, [R0]              ; Write modified NSACR value
ISB                               ; Instruction synchronization barrier
```

Enabling the FPU in CPACR (via CMSIS)

https://github.com/ARM-software/CMSIS_5/blob/develop/Device/ARM/ARMCM33/Source/system_ARMCM33.c

```
76 void SystemInit (void)
77 {
78
79 #if defined (__VTOR_PRESENT) && (__VTOR_PRESENT == 1U)
80     SCB->VTOR = (uint32_t) &(__VECTOR_TABLE[0]);
81 #endif
82
83 #if defined (__FPU_USED) && (__FPU_USED == 1U)
84     SCB->CPACR |= ((3U << 10U*2U) |           /* enable CP10 Full Access */
85                  (3U << 11U*2U) );          /* enable CP11 Full Access */
86 #endif
87
88 #ifndef UNALIGNED_SUPPORT_DISABLE
89     SCB->CCR |= SCB_CCR_UNALIGN_TRP_Msk;
90 #endif
```

```
LDR    R0, =0xE000ED88
LDR    R1, [R0]
ORR    R1, R1, #(0xF << 20)
STR    R1, [R0]
ISB
```

Agenda

- Floating-point Extension overview
- Registers
- Enabling the FPU
- **Floating-point instructions**
- Exceptions

Floating-point instructions (1)

- **Floating-point instructions are of the form:**
`Vop{<c>}{<q>}{.<dt>}{<dest>}, <src1>, <src2>`
- **Where:**
 - <c> and <q> are standard assembler condition and encoding qualifier syntax fields
 - <dt> is a data type specifier, of the form:
 - Floating-point - F16 (half), F32 (single), or F64 (double). F32 can be replaced with F, F64 with D
 - More than one <dt> specifier can be applied to an instruction
- **All Floating-point operations start with a V prefix**
- **F64 only supported for register loads and stores on single-precision variants**
- **F16 only supported for conversion instructions**

Floating-point instructions (2)

- **VMOV used to transfer data between core (R_n) and Floating-point (S_n/D_n) registers:**
 - From one R register to one S register
 - From two R registers to two S registers
 - From one S register to one R register
 - From two S registers to two R registers
 - From two R registers to one D register
- **VMOV immediate available**
- **VMRS and VMSR is used to move a core (R_n) register to and from FPSCR**
- **Transfers of values to and from memory, usual instructions, just with V prefix**
 - VLDR, VLDM, VSTR, VSTM, VPUSH, VPOP
- **Lazy load multiple and lazy store multiple : VLLDM, VLSTM**
 - only available when Security Extension is implemented

Floating-point instructions (3)

- Standard arithmetic instructions: `VADD`, `VSUB`, `VMUL`, `VDIV`
- Fused multiply accumulate, fused multiply subtract: `VMLA`, `VMLS`
- Multiply and negate: `VNMUL`
- Fused negate multiply accumulate, fused negate multiply subtract - negates the multiply before addition/subtraction: `VNMLA`, `VNMLS`
- Negate: `VNEG`
- Select minimum or maximum: `VMIN`, `VMAX`
- Square root (based on method of least squares): `VSQRT`
- Half-precision conversions: `VCVTB` (bottom) `VCVTT` (top)
- Floating-point compares: `VCMP`, `VCMPE`
- Conditional selection selects from one of two registers depending on `FPSCR` condition codes: `VSEL`

Agenda

- Floating-point Extension overview
- Registers
- Enabling the FPU
- Floating-point instructions
- **Exceptions**

Floating-point exceptions

- **Floating-point exceptions are not hardware exceptions**
- **Five IEEE 754-2008 defined Floating-point exceptions**
 - Inexact (IXC)
 - Underflow (UFC)
 - Overflow (OFC)
 - Division by Zero (DZC)
 - Invalid Operation (IOC)
- **One Arm-specific Floating-point exception**
 - Input Denormal (IDC)
 - Can only be set in Flush-to-zero mode (FPSCR.FZ set to 1)
 - Half-precision numbers are exempt from Flush-to-zero mode
- **Once set, stay set until cleared by software (the 'C' stands for 'Cumulative')**

Exception handling

- **Four possible modes for stacking FP context information on taking an exception:**
- **Do not stack any FP context**
 - The processor stacks only a Basic frame
- **Stack an Extended frame**
 - Containing the Basic frame and the FP state information
 - This preserves the Floating-point state required by the AAPCS
- **Reserve space on the stack for an Extended frame, but write only the Basic frame information**
 - No data is pushed into the stack locations reserved for $S0-S15$ and the $FPSCR$ value
 - This is an FP lazy context save
- **Stack entire Floating-point context**
 - All of $S0-S31$ and the $FPSCR$ is stacked
- **The $ASPEN$, $LSPEN$ and TS bits in $FPSCR$, and $CONTROL.FPCA$ determine which action is taken, along with the security state at the time of the exception, and the security of the exception being taken ...**

Floating-point context control

- **Bit 2 of CONTROL register, FPCA, indicates Floating-point Extension is active**
 - Reset value is 0
 - Set automatically using FPCCR.ASPEN (see below)
 - When set, hardware preserves Floating-point context on exception entry, if necessary, and restores on exception return
- **On exception entry, EXC_RETURN[4] is set to inverse of CONTROL.FPCA**
 - EXC_RETURN value of 0xFFFFFEx indicates extended stack frame
- **Floating-point Context Control Register (FPCCR)**
 - Bit 31 **ASPEN** – when set (default), successful execution of any Floating-point instruction sets CONTROL.FPCA to 1
 - Bit 30 **LSPEN** – when set (default), enables Extended stack frame when set
 - Only accessible by privileged software

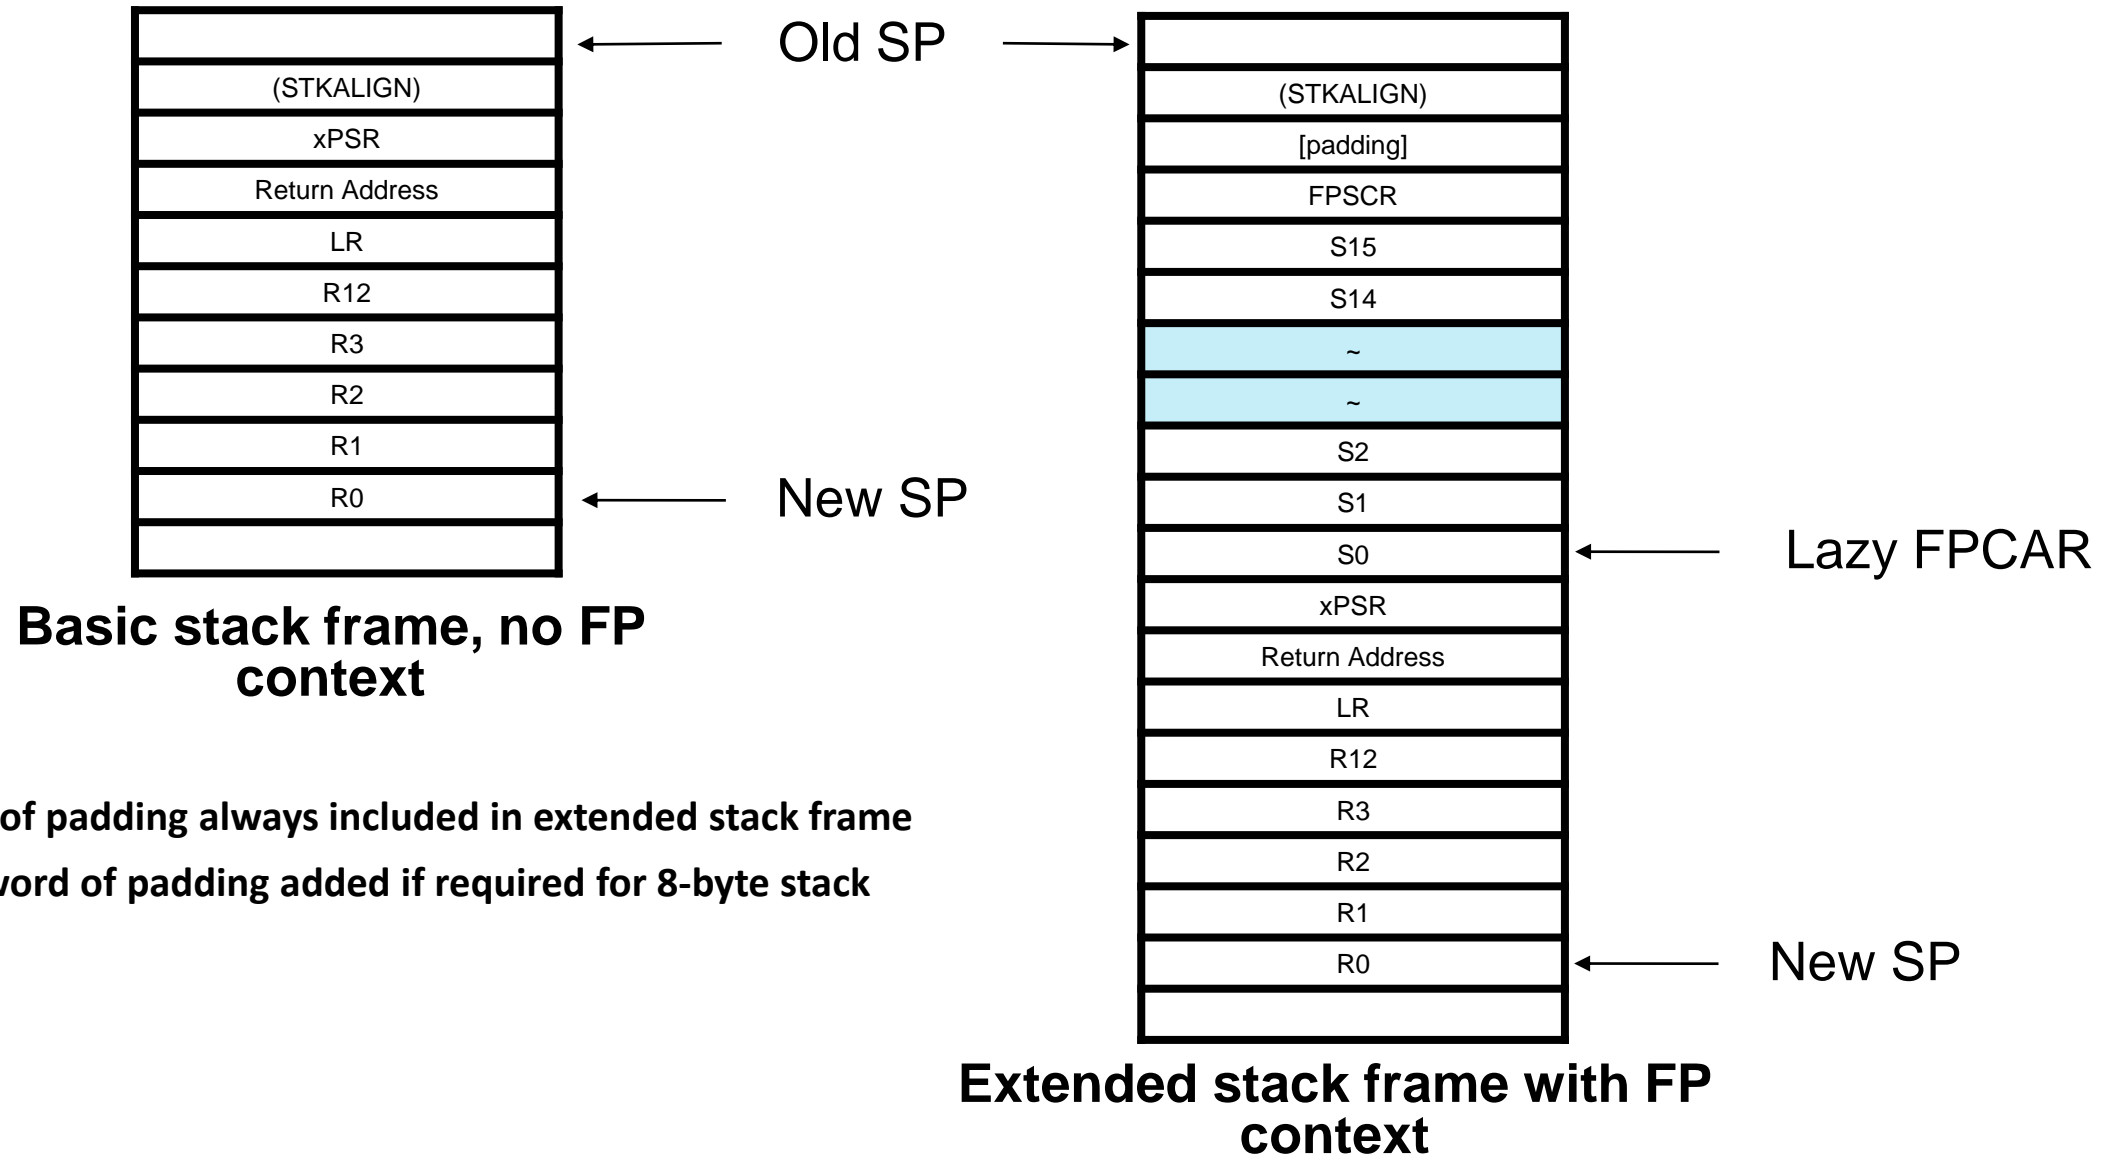
ASPEN=0 LSPEN=0	ASPEN=0 LSPEN=1	ASPEN=1 LSPEN=0	ASPEN=1 LSPEN=1
<ul style="list-style-type: none"> Stack Basic frame 	<ul style="list-style-type: none"> Invalid configuration 	<ul style="list-style-type: none"> Automatic FP stacking Stack Extended frame 	<ul style="list-style-type: none"> Reserve space for Extended stack frame Stack Basic frame Set LPPACT to 1 Stack FP registers only if used in interrupt handler

ASPEN
Automatic State Preservation Enable

LSPEN
Lazy State Preservation Enable

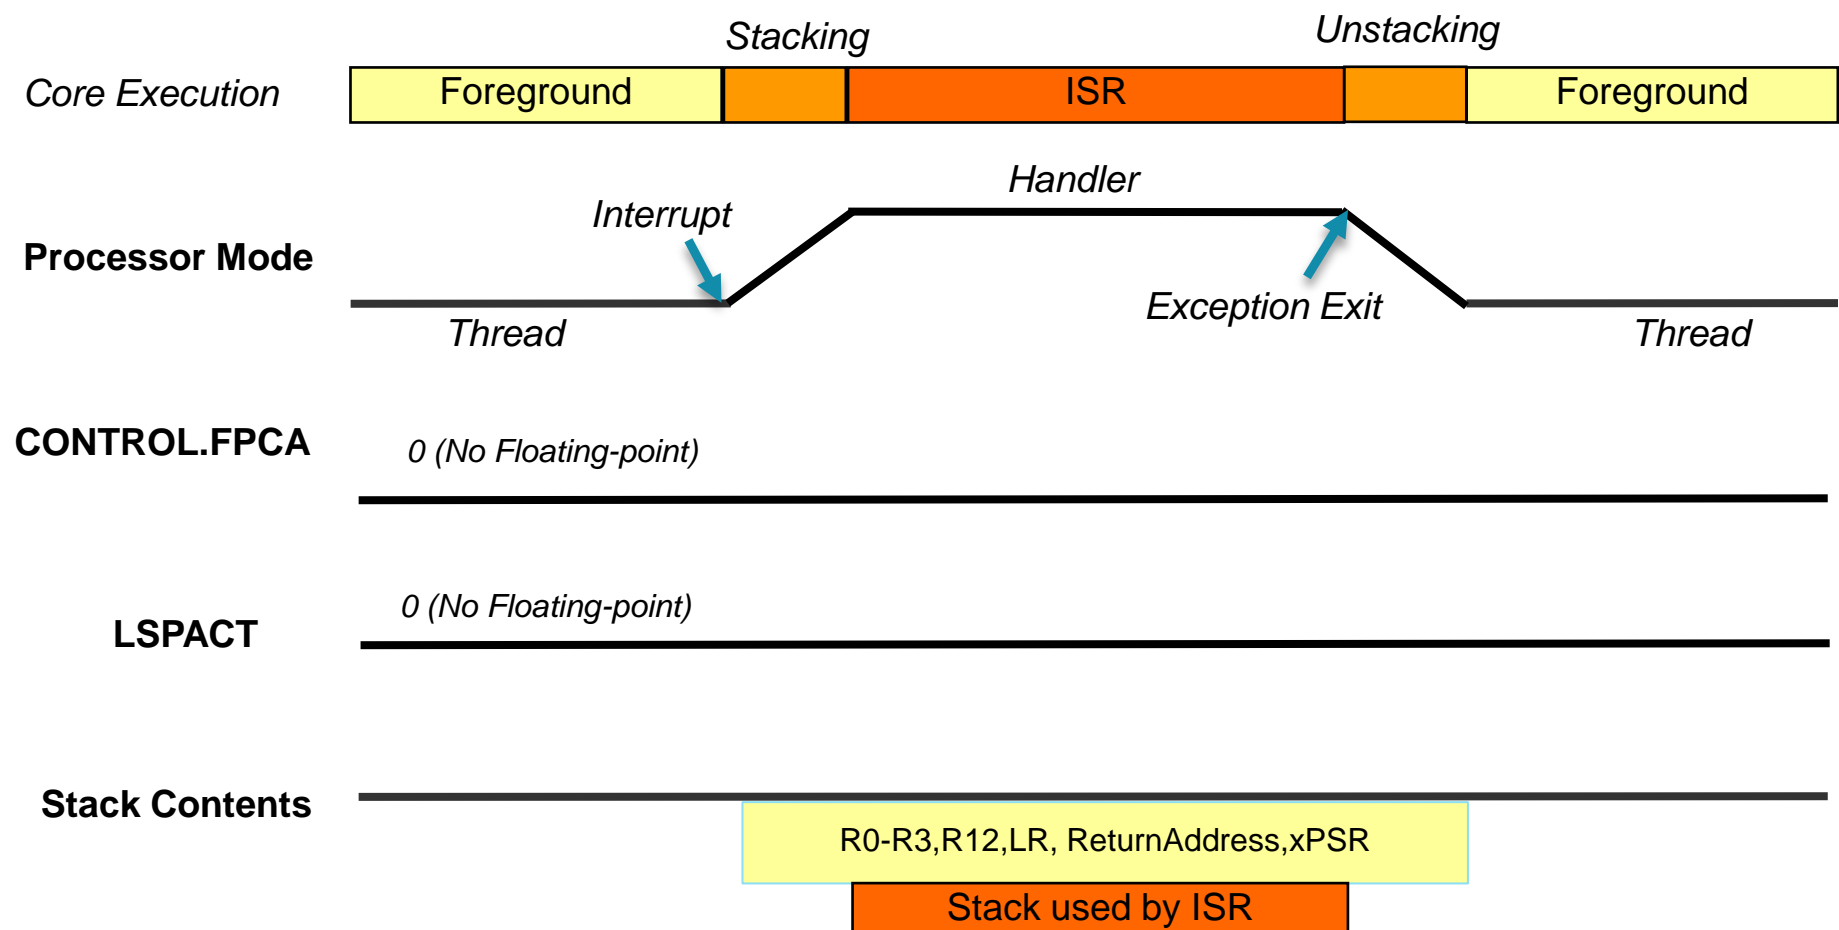
LPPACT
Lazy State Preservation Active

Basic versus Extended frame



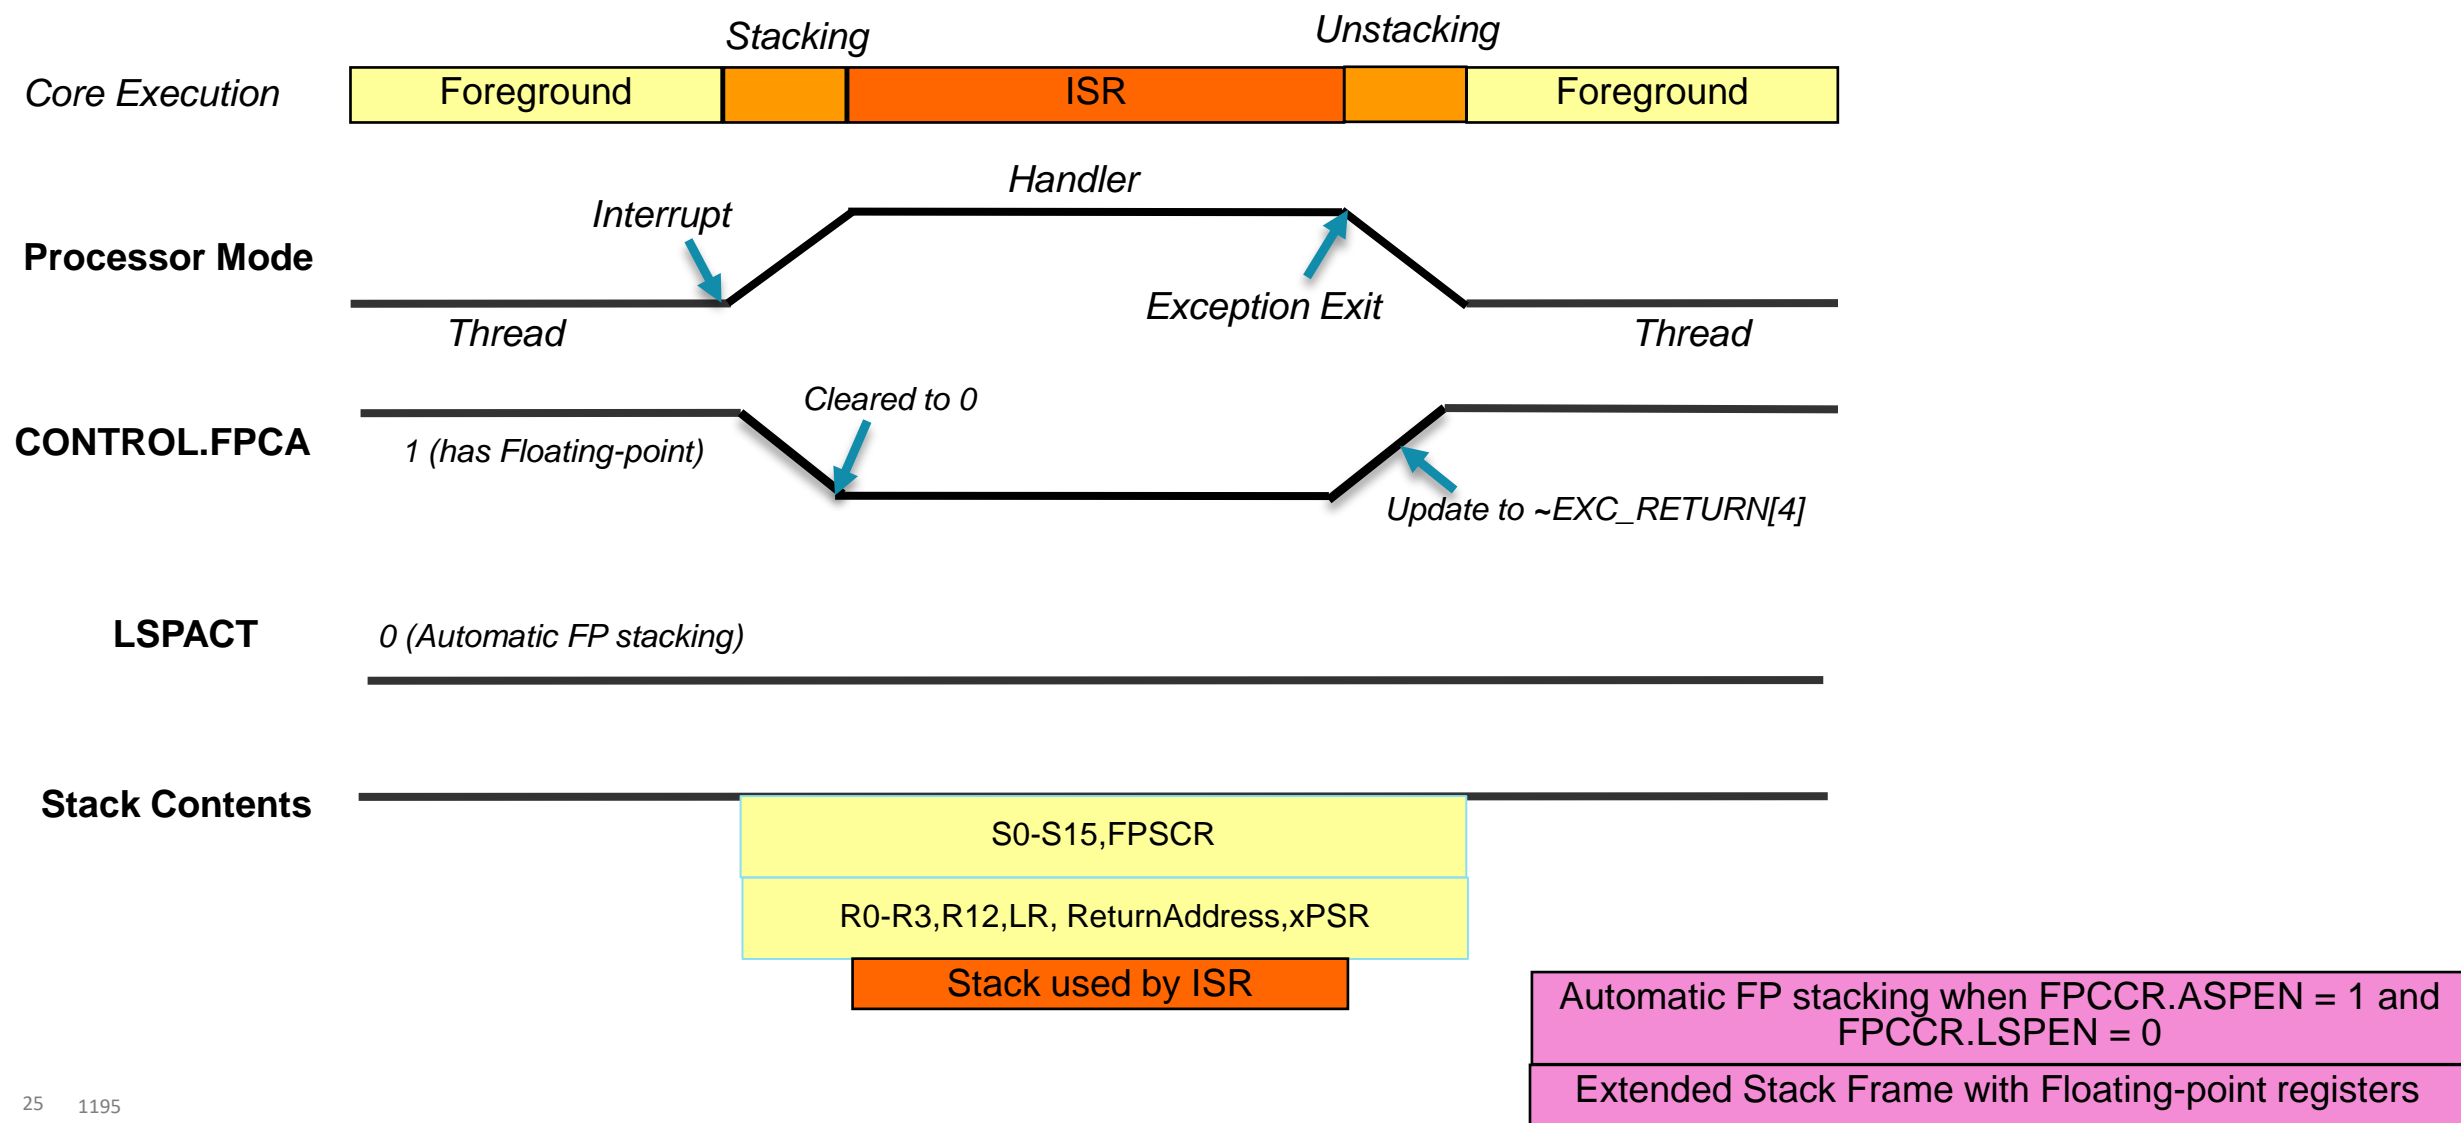
- One word of padding always included in extended stack frame
- Optional word of padding added if required for 8-byte stack alignment

Example 1 – No Floating-point context



Basic Integer stacking when FPCCR.ASPEN = 0 and FPCCR.LSPEN = 0

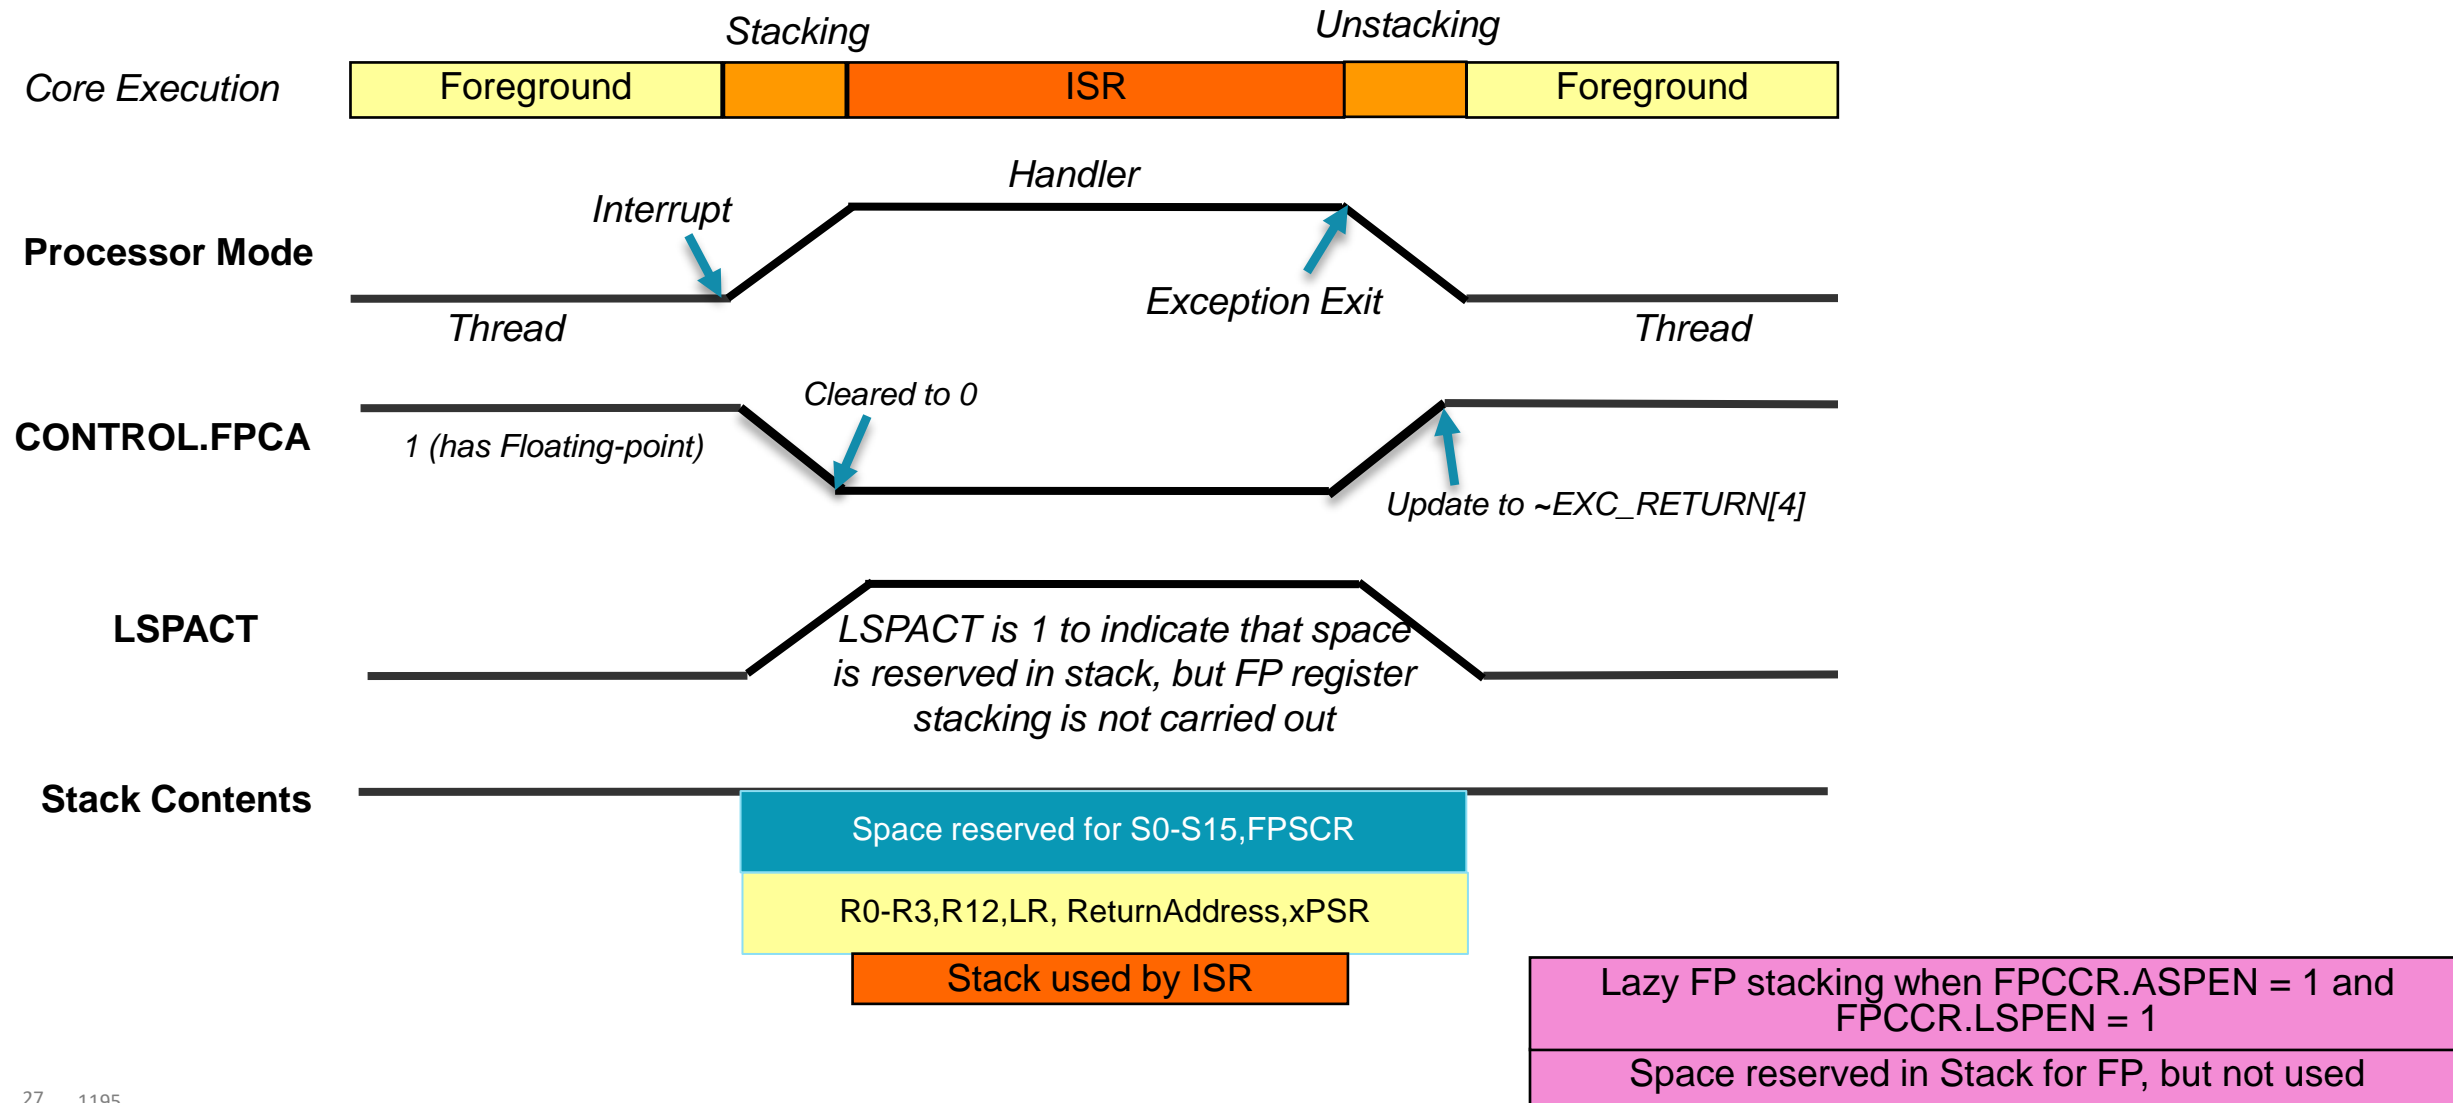
Example:2 With Floating-point context (Automatic)



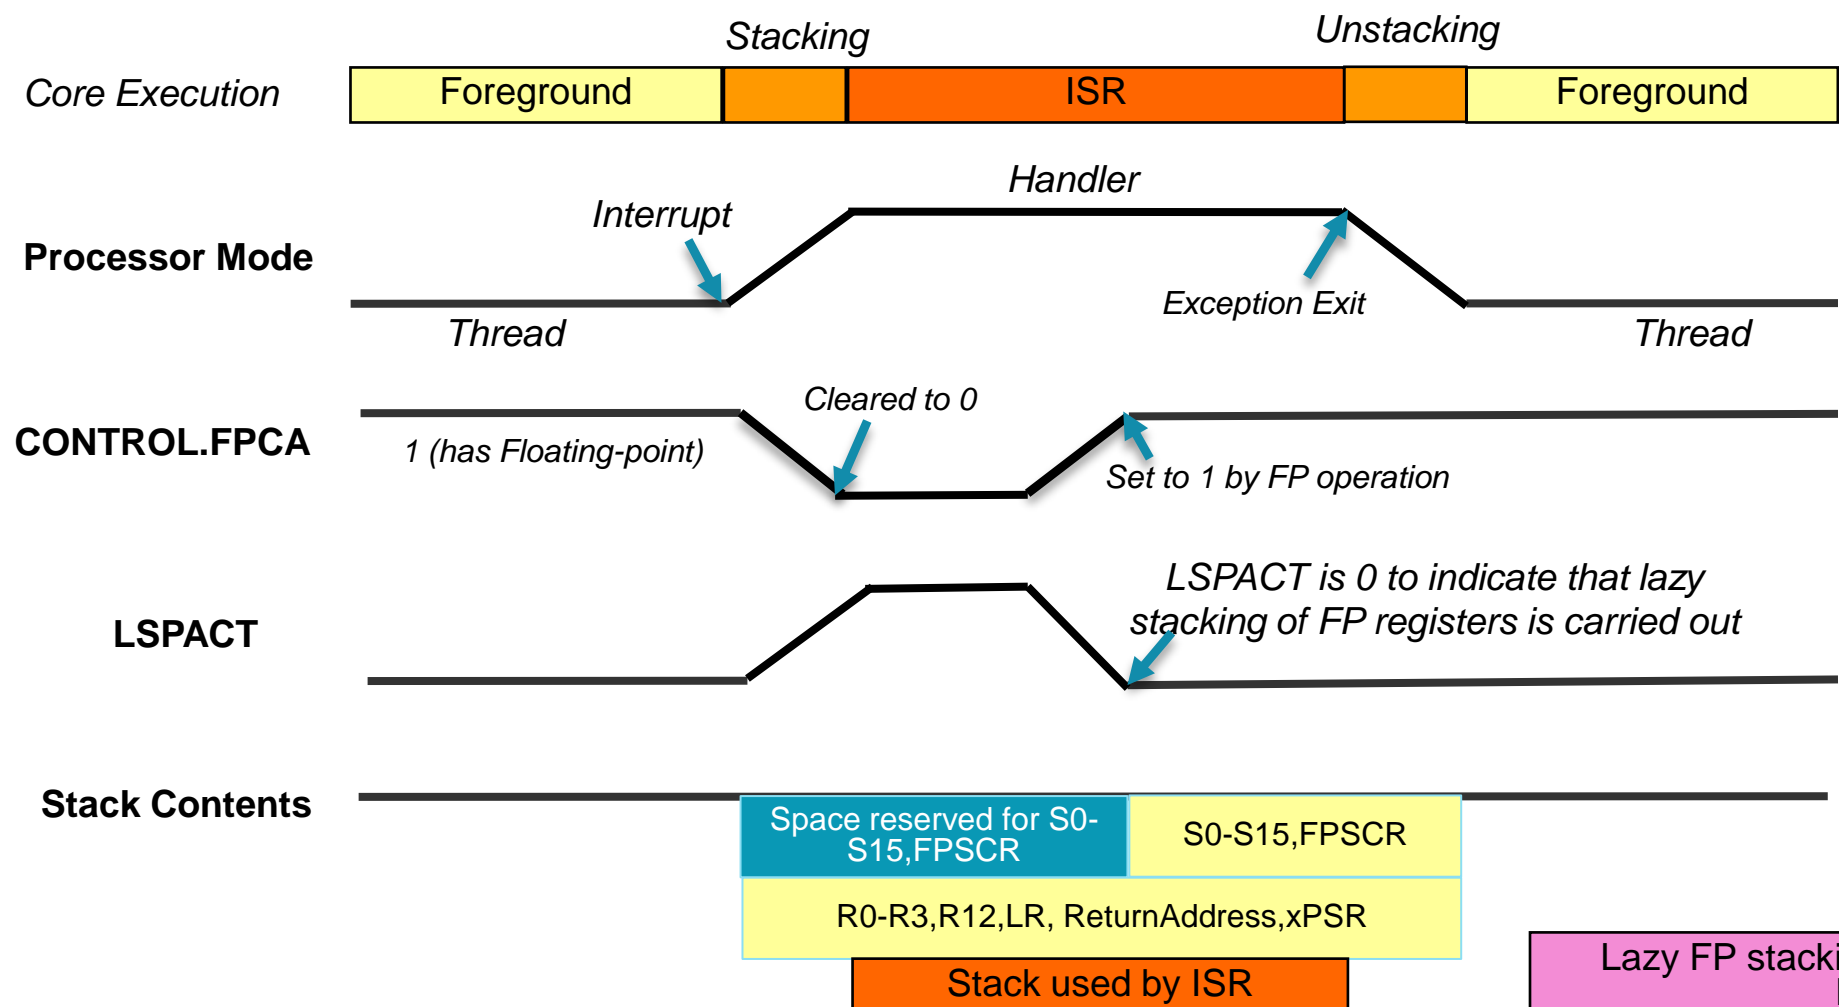
Lazy context save

- Pushing 17 words onto the stack is expensive
- Exception handlers or tasks do not have to use Floating-point registers
- If `FPCCR.LSPEN` is set, then on exception, the stack pointer is incremented as if the Floating-point context is saved, but the push is not performed
- Exception handler can then choose whether to stack or not
- Floating-point Context Address Register (FPCAR) holds the address location of the unpopulated but allocated space in the FP stack frame when lazy stacking is in use
 - Banked between security state
 - Non-secure version `FPCAR_NS`, if Security Extension present
- **IMPORTANT** – `FPCCR.LSPEN` cannot be set while `CONTROL.FPCA` is set
 - Therefore, lazy context switch must be enabled before user application Floating-point code starts executing

Example:3 – Floating-point context in Foreground (Lazy)



Example:4 – Floating-point context in ISR (Lazy)



Lazy FP stacking when FPCCR.ASPEN = 1 and FPCCR.LSPEN = 1

Lazy stacking carried out with space reserved in Stack

Interaction with the Security Extension

- **When the PE also contains the Security Extension, there is an additional interaction**
- **If the the PE is in the the secure state, there is a Floating-point context active, and a non-secure exception is taken:**
 - All the general purpose registers are stacked
 - All 32 Floating-point registers are stacked
 - The xPSR and FPSCR are stacked
 - There are two alignment words which are reserved on the stack
 - There is an integrity signature stacked
 - The return address is stacked
- **This makes the stack frame 51 words long, which is extremely expensive**
- **Think very carefully about using Floating-point in the secure state**

References

- **Armv8-M Architecture Reference Manual**
- **Application Note 298 - Cortex-M4(F) Lazy Stacking and Context Switching**
- **“The Definitive guide to Arm Cortex-M23 and Cortex-M33 processors” – By Joseph Yiu**

arm

Thank You

Danke

Gracias

Grazie

谢谢

ありがとう

Asante

Merci

감사합니다

धन्यवाद

Kiitos

شكراً

ধন্যবাদ

תודה